

HONG WANG

AI APPLICATIONS ENGINEER / AGENT SYSTEMS

Suzhou, China | +86 186 2706 1195 | wh_5233@163.com

GitHub: github.com/wanghong5233 | Portfolio: wh5233.me

EDUCATION

Soochow University, Suzhou, China

Sep 2024 – Jun 2027

M.S. in Artificial Intelligence

Wuhan University of Science and Technology, Wuhan, China

Sep 2019 – Jun 2024

B.S. in Software Engineering

EXPERIENCE

Shanghai Qiwan Zhichi Technology Co., Ltd. — AI Full-Stack Engineer

May 2026 – Sep 2026

- Shipped an AI-assisted script analysis and rewriting service and integrated it into the main product, connecting domain context, evidence retrieval, editor actions, and user-reviewed diffs.
- Designed context routing for long-form scripts across document, chapter, and scene views; stabilized retrieval with query construction, multi-source retrieval, reranking, fallback handling, and evidence lookup.
- Cut event-extraction time for a 345K-character script from 6,962 seconds to approximately 60 seconds by combining hierarchical MapReduce, adaptive output budgets, binary window splitting, and gap recovery; coverage checks rejected incomplete runs.
- Made long-running LLM jobs resumable with PostgreSQL-backed task state, atomic job claiming, leases, bounded retries, cancellation, repair jobs, guarded writes, and concurrency locks.

SELECTED PROJECTS

RivalLens — Multi-Agent competitive intelligence system

May 2026 – Jun 2026

- Led the architecture and core implementation of a ByteDance AI Full-Stack Challenge project that placed in the top 30 of approximately 60 teams.
- Turned natural-language research requests into evidence-backed reports with LangGraph agents for planning, research, analysis, writing, and quality review; used typed Pydantic contracts and automated checks to send unsupported claims back for correction.

Pulse — Personal AI Agent platform

Mar 2026 – May 2026

- Built a reusable Agent runtime with task orchestration, checkpoint/resume, safety controls, observability, and integrations with Lark (Feishu) and WeChat.
- Used the runtime to automate job-search workflows and deliver a recurring technology radar through conversational and scheduled execution.

ScholarMind — AI research assistant

Jul 2025 – Mar 2026

- Built a research assistant around hybrid retrieval, multi-step research agents, and a ReAct document editor, with page-level evidence for each citation.
- Combined Multi-Query and HyDE rewriting with BM25/vector retrieval, RRF/MMR fusion, Cross-Encoder reranking, and page/BBox-level evidence anchors.

RESEARCH & PUBLICATIONS

- First author of “Safe Multi-Agent Deep Reinforcement Learning for Privacy-Aware Edge-Device Collaborative DNN Inference,” published in IEEE Transactions on Vehicular Technology (Aug. 2026).
- Second author of “Graph-Enhanced Deep Reinforcement Learning for Share-Aware Multi-Task DNN Inference in Low-Altitude Intelligent Service Networks,” under review at IEEE Transactions on Services Computing.
- Third author of “Privacy-Aware Joint DNN Model Deployment and Partitioning Optimization for Collaborative Edge Inference Services,” published in IEEE Transactions on Services Computing (2025).

TECHNICAL SKILLS

- **Languages & Systems:** Python, TypeScript, FastAPI, SQL, PostgreSQL, Redis, Docker, Linux, Git
- **Agent & LLM:** LangGraph, ReAct, Multi-Agent, RAG, Tool Calling, MCP, Context Engineering, PyTorch, QLoRA, SFT, DPO